




# So Far...

- ▶ Our goal (supervised learning)
  - Naïve Bayesian classifier
  - Linear Regression
  - Logistic Regression
  - SVM
  - Perceptron
  - Neural Network
  - k Nearest Neighbor
  - Decision Tree
  - **Single Classifier**
- ▶ Can the classifiers be combined to achieve better performance?
  - Two heads are better than one



$$\text{EPE}(f) = (\text{bias})^2 + \text{variance} + \text{noise}$$

# Ensemble Methods & Random Forest

**Deng Cai (蔡登)**

College of Computer Science  
Zhejiang University

dengcai@gmail.com





# The Bagging Algorithm


- ▶ The name Bagging came from the abbreviation of *Bootstrap Aggregating* [Breiman, 1996d], implies the two key ingredients of Bagging are bootstrap and aggregation.
  - **Bootstrap sampling:** re-sample  $N$  examples from  $D$  **uniformly with replacement**—can also use arbitrary  $N'$  instead of original  $N$
  - **Aggregating:** adopts the most popular strategies for aggregating the outputs of the base learners, that is, *voting* for classification and *averaging* for regression.



# The Bagging Algorithm

Pocket is a modified Perceptron Learning Algorithm.

Bagging 25 Pocket Models. The max iteration for each Pocket Algorithm is 1000.



$$T_{\text{POCKET}} = 1000; T_{\text{BAG}} = 25$$

- ▶ very diverse  $g_t$  from bagging
- ▶ proper **non-linear** boundary after aggregating binary classifiers
- ▶ bagging works reasonably well if **base algorithm sensitive to data randomness**



# Why Bagging Works

- ▶ Let  $f$  denote the ground-truth function and  $h(\mathbf{x})$  denote a learner trained from the bootstrap distribution  $D_{bs}$ . The aggregated learner generated by Bagging is

$$H(\mathbf{x}) = E_{D_{bs}}[h(\mathbf{x})]$$



- ▶ With simple algebra and the inequality  $(E[X])^2 \leq E[X^2]$ , we have

$$(f - H(\mathbf{x}))^2 \leq E_{D_{bs}}[(f - h(\mathbf{x}))^2]$$

- ▶ Thus, by integrating both sides over the distribution, we can get that the mean-squared error of  $H(\mathbf{x})$  is smaller than that of  $h(\mathbf{x})$  averaged over the bootstrap sampling distribution



# Why Bagging Works: Bias & Variance

$$\{E_D(f(\mathbf{x}; D)) - E(y|\mathbf{x})\}^2 + E_D\{[f(\mathbf{x}; D) - E_D(f(\mathbf{x}; D))]^2\}$$

- ▶ Let  $f$  denote the ground-truth function and  $h(\mathbf{x})$  denote a learner trained from the bootstrap distribution  $D_{bs}$ . The aggregated learner generated by Bagging is

$$H(\mathbf{x}) = E_{D_{bs}}[h(\mathbf{x})]$$

- ▶ Bias:  $E_D[h(\mathbf{x})] - f \approx H(\mathbf{x}) - f$

- ▶ Variance:

- Assume there are  $T$  *bs* samples, if data is *i. i. d.* and each variance is  $\sigma^2$ . The ensemble variance is  $\frac{1}{T}\sigma^2$

$$H(\mathbf{x}) = E_{D_{bs}}[h(\mathbf{x})] = \frac{1}{T} \sum_{i=1}^T h_i(\mathbf{x})$$



$$\text{Var}(H(\mathbf{x})) = \text{Var}\left(\frac{1}{T} \sum_{i=1}^T h_i(\mathbf{x})\right) = \frac{1}{T^2} \text{Var}\left(\sum_{i=1}^T h_i(\mathbf{x})\right) = \frac{1}{T^2} \sum_{i=1}^T \text{Var}(h_i(\mathbf{x})) = \frac{1}{T} \sigma^2$$

- In reality, Bootstrap Sampled data is *i. d.* (not necessarily independent) with correlation  $\rho$

$$\rho\sigma^2 + \frac{1-\rho}{T}\sigma^2$$




# Why Bagging Works



**Why?**

If Bootstrap Sampled data is  
*i. d.* with correlation  $\rho$ :

$$\rho\sigma^2 + \frac{1-\rho}{T}\sigma^2$$







# Why Bagging Works

- ▶ Performance improvement brought by Bagging is large when the base learner is unstable (large variance).
  - The base learner should be aware of the little change, sometimes overfitting is allowable.
- ▶ Thus, Friedman and Hall [2007] concluded that Bagging can reduce the variance of higher-order components, yet not affect the linear components. This implies that Bagging is better applied with highly nonlinear learners.

*J. H. Friedman and P. Hall. On bagging and nonlinear estimation. Journal of Statistical Planning and Inference, 137(3):669–683, 2007.*



# Which Classifier is a good choice for base learner?

- Naïve Bayesian classifier
- Perceptron
- Linear Regression
- Logistic Regression
- SVM
- Neural Network
- k Nearest Neighbor
- Decision Tree

Model with **low bias** benefits from bagging

## ► Decision Tree!

- Non-linear classifier
- Easy to use and interpret
- Can perfectly fit to any training data (overfitting) with high test error. (**zero bias, high variance**)




# Random Forest

- ▶ Random Forest, which uses **decision tree as basic unit** in bagging, is an ensemble model.
- ▶ Why Random Forest?
  - Advantages of Decision Tree:
    - Handling missing value
    - Robust to outliers in input space
    - Fast
    - Good Interpretability
  - Limitations of Decision Tree:
    - Low accuracy, low bias with high variance and easy to overfit
  - Ensemble to maintain advantages while increasing accuracy




# Random Forest


- ▶ Bagging sample is a random process.
  - Bootstrap sample process: random sampling the given  $N$  samples with replacement
  - Repeat  $T$  times and create new  $T$  training sets
  - Get  $T$  models
- ▶ Aggregating all models will fuzzy up the decision boundary, which help reduce the variance, and prevent the *one man rule* danger



Data Set





Decision Tree



Random Forest



# How to increase the variation of each tree?





# Random Forest

- ▶ The main difference between the method of RF and Bagging is that RF incorporate **randomized feature selection at each split step**.
- ▶ At each split step : Get feature subsets  $f$  from the whole  $F$ 
  - More efficient in building trees
    - each time we only pick the best feature from size( $f$ ) rather than size( $F$ ).
      - We often let  $\text{size}(f) = \sqrt{k}$  in classification and  $k/3$  in regression
  - Each tree is not as good as DT, but work better when they are combined.



# Construct Random Forest

- ▶ Let  $N$  be the number of trees and  $K$  be the feature subset size.
- ▶ For each  $N$  iterations: (building a tree in forest)
  - 1. Select a new bootstrap sample from original training set
  - 2. Growing a tree...
  - 3. At each internal node, randomly select  $K$  features from ALL features and then, determine the best split in ONLY the  $K$  features.
  - 4. Do not pruning
  - 5. Until Test Error never decrease (here means Validation error in RF)
- ▶ At last, overall prediction as the average( or vote) from  $N$  trees.



## 3 Rules of RF makes the learners more **Diverse**.

- ▶ Random forest need basic learner **aware the little change**, sometimes overfit is allowable.
- ▶ Each time, basic learner doesn't learn from all data, but from **Random** bootstrap sampled data.
- ▶ Basic learner doesn't use all features, but **Random** select some features.





# Other ways to generate base learner?

- ▶ The base learner should have low bias? **Not necessary**
- ▶ The base learner should have high variance? **What does this mean?**
  - Different base learners should be different.
  - Different base learners should have different result.


$$\min_f \left\{ \sum_{i=1}^n \ell(f, \mathbf{x}_i, y_i) + \lambda R(f) \right\}$$

$$\min_f \left\{ \sum_{i=1}^n w_i \ell(f, \mathbf{x}_i, y_i) + \lambda R(f) \right\}$$

- ▶ Boosting
  - Training base learners by assigning different weights to the samples




# Boosting






# Boosting






# Boosting






# Boosting






# Boosting





# Boosting





# Observation

- ▶ In order to come up an algorithm, we need to answer two underlying questions:
  - Question 1: How to change the weights of samples so that misclassified samples get more weight
  - Question 2: How to combine base learners in final phase





# AdaBoost

- ▶ 1. Initialize the data weighting coefficients  $\mathbf{w}$  by setting  $w_n^{(1)} = 1/N$  for  $n = 1, \dots, N$ .
- ▶ 2. For  $m = 1, \dots, M$ :
  - (a) Fit a classifier  $y^{(m)}(\mathbf{x})$  to the training data by minimizing the weighted error function

$$J_m = \sum_{n=1}^N w_n^{(m)} I(y^{(m)}(\mathbf{x}_n) \neq t_n)$$



# AdaBoost

(b) Evaluate the Errorrate:

$$\epsilon_m = \frac{\sum_{n=1}^N w_n^{(m)} I(y^{(m)}(x_n) \neq t_n)}{\sum_{n=1}^N w_n^{(m)}}$$

and then use this to evaluate

$$\alpha_m = \ln \frac{1 - \epsilon_m}{\epsilon_m}$$

Log odds, smaller the error rates, bigger this value.

(c) the data weighting coefficients

$$w_n^{(m+1)} = w_n^{(m)} \exp\{\alpha_m I(y^{(m)}(x_n) \neq t_n)\} = \begin{cases} w_n^{(m)} \frac{1 - \epsilon_m}{\epsilon_m}, & \text{if } y^{(m)} \text{ makes error} \\ w_n^{(m)}, & \text{otherwise} \end{cases}$$

Answer to question 1, How to change the weights of samples so that misclassified samples get more weight.



# AdaBoost

- ▶ 3. Make predictions using the final model, which is given by

$$Y_M(\mathbf{x}) = \text{sign} \left( \sum_{m=1}^M \alpha_m y^{(m)}(\mathbf{x}) \right)$$

Question 2: How to  
Combine base learners in  
final phase

weighting coefficients  $\alpha_m$  give greater weight to the more accurate classifiers when computing the overall output



# Insights Behind Adaboost

- ▶ In each round, the algorithm try to get a different base learner, so that model diversity is achieved



# Re-weight for More Diverse Base Learner

$$g_m \longleftarrow \min_{h \in H} \left( \sum_{n=1}^N w_n^{(m)} I(y_n^{(m)} \neq t_n) \right)$$

$$g_{m+1} \longleftarrow \min_{h \in H} \left( \sum_{n=1}^N w_n^{(m+1)} I(y_n^{(m+1)} \neq t_n) \right)$$

if  $g_m$  'not good' for  $w^{(m+1)} \Rightarrow g_{m+1}$  diverse from  $g_m$

**Idea: construct  $w^{m+1}$  to make  $g_m$  random-like**

$$\frac{\sum_{n=1}^N w_n^{(m+1)} I(y_n^{(m)} \neq t_n)}{\sum_{n=1}^N w_n^{(m+1)}} = \frac{1}{2} \quad \frac{\sum_{n=1}^N w_n^{(m+1)} I(y_n^{(m)} \neq t_n)}{\sum_{n=1}^N w_n^{(m+1)} I(y_n^{(m)} \neq t_n) + \sum_{n=1}^N w_n^{(m+1)} I(y_n^{(m)} = t_n)} = \frac{1}{2}$$

Solve this equation will give us

$$w_n^{(m+1)} = w_n^{(m)} \exp \left\{ -\frac{1}{2} \alpha_m t_n y^{(m)}(x_n) \right\}$$

$$w_n^{(m+1)} = w_n^{(m)} \sqrt{\frac{1 - \epsilon_m}{\epsilon_m}}$$

$$w_n^{(m+1)} = w_n^{(m)} \sqrt{\frac{\epsilon_m}{1 - \epsilon_m}}$$

$$w_n^{(m+1)} = w_n^{(m)} \frac{1 - \epsilon_m}{\epsilon_m}$$

which is the same as update equation in Adaboost Algorithm



# Insights Behind Adaboost

- ▶ In each round, the algorithm try to get a different base learner, so that model diversity is achieved
- ▶ Adaboost can be see as a **sequential** optimization process of an additive model under exponential error



# Optimization of Exponential Loss

- Consider the exponential error function defined by

$$E = \sum_{n=1}^N \exp\{-t_n f_m(\mathbf{x}_n)\}$$

where  $f_m(\mathbf{x})$  is a classifier defined in terms of a linear combination of base classifiers  $y_l(\mathbf{x})$  of the form

$$f_m(\mathbf{x}) = \frac{1}{2} \sum_{l=1}^m \alpha_l y_l(\mathbf{x})$$

and  $t_n \in \{-1, 1\}$  are the training set target values. Our goal is to minimize  $E$  with respect to both the weighting coefficients  $\alpha_l$  and the parameters of the base classifiers  $y_l(\mathbf{x})$ .

- Instead of doing a global error function minimization, however, we shall suppose that the base classifiers  $y_1(\mathbf{x}), \dots, y_{m-1}(\mathbf{x})$  are fixed, as are their coefficients  $\alpha_1, \dots, \alpha_{m-1}$ , and so we are minimizing only with respect to  $\alpha_m$  and  $y_m(\mathbf{x})$ .



# Optimization of Exponential Loss

$$E = \sum_{n=1}^N \exp\{-t_n f_m(\mathbf{x}_n)\} \quad f_m(\mathbf{x}) = \frac{1}{2} \sum_{l=1}^m \alpha_l y_l(\mathbf{x})$$

- ▶ the base classifiers  $y_1(\mathbf{x}), \dots, y_{m-1}(\mathbf{x})$  are fixed, as are their coefficients  $\alpha_1, \dots, \alpha_{m-1}$ ,

$$\begin{aligned} &= \sum_{n=1}^N \exp\left\{-t_n f_{m-1}(\mathbf{x}_n) - \frac{1}{2} t_n \alpha_m y_m(\mathbf{x}_n)\right\} \\ &= \sum_{n=1}^N \exp\{-t_n f_{m-1}(\mathbf{x}_n)\} \exp\left\{-\frac{1}{2} t_n \alpha_m y_m(\mathbf{x}_n)\right\} \\ &= \sum_{n=1}^N w_n^{(m)} \exp\left\{-\frac{1}{2} t_n \alpha_m y_m(\mathbf{x}_n)\right\} \end{aligned}$$

$$w_n^{(m+1)} = w_n^{(m)} \exp\left\{-\frac{1}{2} t_n \alpha_m y_m(\mathbf{x}_n)\right\} \quad t_n y_m(\mathbf{x}_n) = 1 - 2I(y^{(m)}(\mathbf{x}_n) \neq t_n)$$

$$w_n^{(m+1)} = w_n^{(m)} \exp\left\{-\frac{\alpha_m}{2}\right\} \exp\{\alpha_m I(y^{(m)}(\mathbf{x}_n) \neq t_n)\}$$





# Optimization of Exponential Loss

- If we denote by  $T_m$  the set of data points that are correctly classified by  $y_m(x)$ , and if we denote the remaining misclassified points by  $M_m$ , then we can in turn rewrite the error function in the form.

$$\begin{aligned} E &= e^{-\frac{\alpha_m}{2}} \sum_{n \in T_m} w_n^{(m)} + e^{\frac{\alpha_m}{2}} \sum_{n \in M_m} w_n^{(m)} \\ &= \left( e^{\frac{\alpha_m}{2}} - e^{-\frac{\alpha_m}{2}} \right) \sum_{n=1}^N w_n^{(m)} I(y_m(x_n) \neq t_n) + e^{-\frac{\alpha_m}{2}} \sum_{n=1}^N w_n^{(m)} \end{aligned}$$

$$J_m = \sum_{n=1}^N w_n^{(m)} I(y^{(m)}(x_n) \neq t_n)$$

When we minimize this with respect to  $y_m(x)$ , we see that the second term is constant, and so this is equivalent to the procedure in Adaboost Algorithm 2-(a) and we get  $\varepsilon_m$  which is the same as in Adaboost. Similarly, minimizing with respect to  $\alpha_m$ , we get the  $\alpha_m$  the same as  $\alpha_m$  in Adaboost.



# Loss functions for boosting


- ▶ The exponential loss function that is minimized by the AdaBoost algorithm is

$$E = \sum_{n=1}^N \exp\{-t_n f_m(x_n)\}$$

- ▶ Why these loss function?
- ▶ What properties do this loss function have?
- ▶ Can we choose other loss functions?
- ▶ Those questions opens the door to a wide range of boosting-like algorithms, including multiclass extensions, by altering the choice of loss function. It also motivates the extension to regression problems



# Comparing to Other Loss Function



- ▶ Comparing exponential loss, logistic loss, and hinge loss
- ▶ Error function grows exponentially with  $|ty(x)|$ , it penalizes large negative values of  $ty(x)$ , thus be much **less robust** to outliers or misclassified data



# Summary

- ▶ They are two different Ensemble Paradigms
  - Boosting is **sequential ensemble methods**, where the base learners are generated sequentially.
  - Bagging is **parallel ensemble methods**: where the base learners are generated in parallel.
  - Boosting exploit the *dependence* between the base learners, since the overall performance can be boosted in a residual-decreasing way.
  - Bagging exploit the *independence* between the base learners, since the error can be reduced dramatically by combining independent base learners.